


## ASSIGNMENT-1

### PRIYACHANDANA KODATI

#### 1) Two concealed layers were employed. Examine the effects of implementing one or three hidden layers on test accuracy and validation.

Using the Keras framework, this code builds a neural network model for binary sentiment categorization on the IMDB dataset. The two hidden layers in the model architecture, each with 16 neurons, employ the ReLU activation function, whereas the single output layer uses a sigmoid activation function. The model is put together using the Adam optimizer, with accuracy acting as a measure and binary crossentropy as the loss function. A training set and a validation set of training data are used to train the model across 20 epochs with a batch size of 512. The model exhibits good performance during training, with high accuracy starting at roughly 96.7% and rising to nearly 100% by the end of the epochs. The precision of validation also demonstrates strong performance. Historical data, including metrics for accuracy and loss for training and validation sets, is kept for further research.

```
Epoch 1/20
30/30 — 5s 86ms/step - accuracy: 0.6271 - loss: 0.6361 - val_accuracy: 0.8503 - val_loss: 0.4102
Epoch 2/20
30/30 — 3s 40ms/step - accuracy: 0.8992 - loss: 0.3196 - val_accuracy: 0.8869 - val_loss: 0.2904
Epoch 3/20
30/30 — 1s 24ms/step - accuracy: 0.9406 - loss: 0.1831 - val_accuracy: 0.8863 - val_loss: 0.2842
Epoch 4/20
30/30 — 1s 21ms/step - accuracy: 0.9642 - loss: 0.1241 - val_accuracy: 0.8839 - val_loss: 0.3028
Epoch 5/20
30/30 — 1s 20ms/step - accuracy: 0.9777 - loss: 0.0879 - val_accuracy: 0.8813 - val_loss: 0.3314
Epoch 6/20
30/30 — 1s 20ms/step - accuracy: 0.9896 - loss: 0.0549 - val_accuracy: 0.8786 - val_loss: 0.3770
Epoch 7/20
30/30 — 1s 21ms/step - accuracy: 0.9930 - loss: 0.0404 - val_accuracy: 0.8733 - val_loss: 0.4288
Epoch 8/20
30/30 — 1s 21ms/step - accuracy: 0.9960 - loss: 0.0258 - val_accuracy: 0.8737 - val_loss: 0.4739
Epoch 9/20
30/30 — 1s 20ms/step - accuracy: 0.9991 - loss: 0.0152 - val_accuracy: 0.8717 - val_loss: 0.5165
Epoch 10/20
30/30 — 1s 20ms/step - accuracy: 0.9994 - loss: 0.0103 - val_accuracy: 0.8705 - val_loss: 0.5615
Epoch 11/20
30/30 — 1s 21ms/step - accuracy: 0.9999 - loss: 0.0069 - val_accuracy: 0.8702 - val_loss: 0.5923
Epoch 12/20
30/30 — 1s 23ms/step - accuracy: 0.9999 - loss: 0.0048 - val_accuracy: 0.8695 - val_loss: 0.6234
Epoch 13/20
30/30 — 1s 20ms/step - accuracy: 1.0000 - loss: 0.0037 - val_accuracy: 0.8687 - val_loss: 0.6512
Epoch 14/20
30/30 — 1s 20ms/step - accuracy: 1.0000 - loss: 0.0029 - val_accuracy: 0.8686 - val_loss: 0.6710
Epoch 15/20
30/30 — 1s 37ms/step - accuracy: 0.9999 - loss: 0.0025 - val_accuracy: 0.8693 - val_loss: 0.6921
Epoch 16/20
30/30 — 1s 26ms/step - accuracy: 0.9999 - loss: 0.0021 - val_accuracy: 0.8684 - val_loss: 0.7125
Epoch 17/20
30/30 — 1s 22ms/step - accuracy: 0.9999 - loss: 0.0019 - val_accuracy: 0.8683 - val_loss: 0.7309
Epoch 18/20
30/30 — 1s 31ms/step - accuracy: 1.0000 - loss: 0.0015 - val_accuracy: 0.8673 - val_loss: 0.7472
Epoch 19/20
30/30 — 1s 21ms/step - accuracy: 0.9999 - loss: 0.0014 - val_accuracy: 0.8680 - val_loss: 0.7652
Epoch 20/20
30/30 — 1s 20ms/step - accuracy: 1.0000 - loss: 0.0011 - val_accuracy: 0.8677 - val_loss: 0.7773
```



```
51] results2=model12.evaluate(x_test,y_test)
```

782/782 — 2s 2ms/step - accuracy: 0.8542 - loss: 0.8519

## 2) Consider utilizing layers of 32, 64, and so on hidden units, or fewer or more hidden units.

With two hidden layers of sixteen neurons each, one output layer with a sigmoid activation function, and a binary classification using Keras, this block of code constructs a neural network model. After the model is constructed using the Adam optimizer with binary crossentropy loss, it is trained on the training dataset for 4 epochs with a batch size of 512. Over the course of training, the model's accuracy rises dramatically, from 73.3% in the first epoch to around 95.3% by the fourth. After training, the model performs reasonably well on unseen data, as evidenced by its accuracy of approximately 87.7% and loss of approximately 0.31 when tested on the test dataset.



```

Epoch 1/20
30/30 — 5s 92ms/step - accuracy: 0.6938 - loss: 0.6152 - val_accuracy: 0.8631 - val_loss: 0.3661
Epoch 2/20
30/30 — 2s 19ms/step - accuracy: 0.9019 - loss: 0.2877 - val_accuracy: 0.8907 - val_loss: 0.2778
Epoch 3/20
30/30 — 1s 28ms/step - accuracy: 0.9455 - loss: 0.1694 - val_accuracy: 0.8876 - val_loss: 0.2837
Epoch 4/20
30/30 — 1s 28ms/step - accuracy: 0.9677 - loss: 0.1141 - val_accuracy: 0.8841 - val_loss: 0.3035
Epoch 5/20
30/30 — 1s 21ms/step - accuracy: 0.9805 - loss: 0.0788 - val_accuracy: 0.8789 - val_loss: 0.3433
Epoch 6/20
30/30 — 1s 22ms/step - accuracy: 0.9881 - loss: 0.0558 - val_accuracy: 0.8789 - val_loss: 0.3697
Epoch 7/20
30/30 — 1s 20ms/step - accuracy: 0.9950 - loss: 0.0352 - val_accuracy: 0.8779 - val_loss: 0.4148
Epoch 8/20
30/30 — 1s 28ms/step - accuracy: 0.9978 - loss: 0.0231 - val_accuracy: 0.8737 - val_loss: 0.4554
Epoch 9/20
30/30 — 1s 27ms/step - accuracy: 0.9996 - loss: 0.0154 - val_accuracy: 0.8747 - val_loss: 0.4921
Epoch 10/20
30/30 — 1s 24ms/step - accuracy: 0.9998 - loss: 0.0101 - val_accuracy: 0.8723 - val_loss: 0.5325
Epoch 11/20
30/30 — 1s 27ms/step - accuracy: 0.9999 - loss: 0.0065 - val_accuracy: 0.8713 - val_loss: 0.5677
Epoch 12/20
30/30 — 1s 20ms/step - accuracy: 0.9996 - loss: 0.0049 - val_accuracy: 0.8713 - val_loss: 0.5988
Epoch 13/20
30/30 — 1s 28ms/step - accuracy: 0.9999 - loss: 0.0033 - val_accuracy: 0.8705 - val_loss: 0.6254
Epoch 14/20
30/30 — 1s 21ms/step - accuracy: 1.0000 - loss: 0.0027 - val_accuracy: 0.8689 - val_loss: 0.6495
Epoch 15/20
30/30 — 1s 23ms/step - accuracy: 0.9999 - loss: 0.0020 - val_accuracy: 0.8686 - val_loss: 0.6712
Epoch 16/20
30/30 — 1s 20ms/step - accuracy: 1.0000 - loss: 0.0016 - val_accuracy: 0.8677 - val_loss: 0.6911
Epoch 17/20
30/30 — 1s 23ms/step - accuracy: 1.0000 - loss: 0.0013 - val_accuracy: 0.8676 - val_loss: 0.7088
Epoch 18/20
30/30 — 1s 23ms/step - accuracy: 1.0000 - loss: 0.0011 - val_accuracy: 0.8673 - val_loss: 0.7250
Epoch 19/20
30/30 — 1s 24ms/step - accuracy: 1.0000 - loss: 0.0010 - val_accuracy: 0.8679 - val_loss: 0.7401
Epoch 20/20
30/30 — 1s 21ms/step - accuracy: 1.0000 - loss: 8.0492e-04 - val_accuracy: 0.8674 - val_loss: 0.7547

```

```
[54] results3=model3.evaluate(x_test,y_test)
```

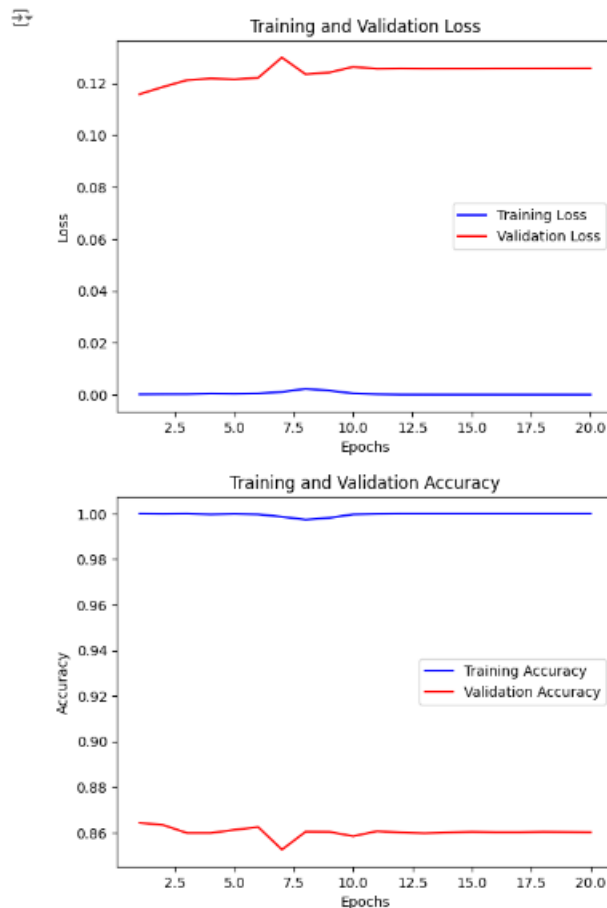
```
782/782 — 3s 3ms/step - accuracy: 0.8567 - loss: 0.8349
```



### 3) Instead of using `binary_crossentropy`, consider utilizing the `mse` loss function.

This code constructs a neural network model (`model3`) that uses ReLU activation and has two hidden layers, each with sixteen neurons. In this definition, `keras` is utilized. The model is assembled using the Adam optimizer, and the loss function is the mean squared error (MSE). It is trained on a subset of the training data for 20 epochs with a batch size of 512, and its performance is monitored with validation data. The algorithm visualizes accuracy over the epochs and training and validation loss using Matplotlib after training. The model yields a test accuracy of roughly 88.0% when evaluated on a test dataset, demonstrating its ability to correctly identify the data.

```
Epoch 1/20
30/30 ————— 4s 79ms/step - accuracy: 1.0000 - loss: 1.0392e-04 - val_accuracy: 0.8643 - val_loss: 0.1158
Epoch 2/20
30/30 ————— 1s 20ms/step - accuracy: 1.0000 - loss: 1.4957e-04 - val_accuracy: 0.8633 - val_loss: 0.1186
Epoch 3/20
30/30 ————— 1s 21ms/step - accuracy: 1.0000 - loss: 1.2279e-04 - val_accuracy: 0.8598 - val_loss: 0.1212
Epoch 4/20
30/30 ————— 1s 23ms/step - accuracy: 0.9999 - loss: 2.1463e-04 - val_accuracy: 0.8598 - val_loss: 0.1218
Epoch 5/20
30/30 ————— 1s 20ms/step - accuracy: 0.9999 - loss: 2.5374e-04 - val_accuracy: 0.8612 - val_loss: 0.1215
Epoch 6/20
30/30 ————— 1s 20ms/step - accuracy: 0.9998 - loss: 2.5967e-04 - val_accuracy: 0.8624 - val_loss: 0.1221
Epoch 7/20
30/30 ————— 1s 30ms/step - accuracy: 0.9994 - loss: 5.3285e-04 - val_accuracy: 0.8525 - val_loss: 0.1299
Epoch 8/20
30/30 ————— 1s 38ms/step - accuracy: 0.9966 - loss: 0.0029 - val_accuracy: 0.8604 - val_loss: 0.1235
Epoch 9/20
30/30 ————— 1s 32ms/step - accuracy: 0.9977 - loss: 0.0020 - val_accuracy: 0.8603 - val_loss: 0.1242
Epoch 10/20
30/30 ————— 1s 21ms/step - accuracy: 0.9997 - loss: 3.9748e-04 - val_accuracy: 0.8584 - val_loss: 0.1263
Epoch 11/20
30/30 ————— 1s 22ms/step - accuracy: 0.9998 - loss: 2.7806e-04 - val_accuracy: 0.8606 - val_loss: 0.1256
Epoch 12/20
30/30 ————— 1s 20ms/step - accuracy: 1.0000 - loss: 3.1734e-05 - val_accuracy: 0.8600 - val_loss: 0.1257
Epoch 13/20
30/30 ————— 1s 21ms/step - accuracy: 1.0000 - loss: 2.2820e-05 - val_accuracy: 0.8597 - val_loss: 0.1256
Epoch 14/20
30/30 ————— 1s 23ms/step - accuracy: 1.0000 - loss: 1.5899e-05 - val_accuracy: 0.8600 - val_loss: 0.1256
Epoch 15/20
30/30 ————— 1s 20ms/step - accuracy: 1.0000 - loss: 1.5185e-05 - val_accuracy: 0.8603 - val_loss: 0.1256
Epoch 16/20
30/30 ————— 1s 21ms/step - accuracy: 1.0000 - loss: 1.3636e-05 - val_accuracy: 0.8601 - val_loss: 0.1257
Epoch 17/20
30/30 ————— 1s 20ms/step - accuracy: 1.0000 - loss: 1.2575e-05 - val_accuracy: 0.8601 - val_loss: 0.1257
Epoch 18/20
30/30 ————— 1s 21ms/step - accuracy: 1.0000 - loss: 1.0872e-05 - val_accuracy: 0.8603 - val_loss: 0.1257
Epoch 19/20
30/30 ————— 1s 21ms/step - accuracy: 1.0000 - loss: 1.0682e-05 - val_accuracy: 0.8602 - val_loss: 0.1257
Epoch 20/20
30/30 ————— 1s 20ms/step - accuracy: 1.0000 - loss: 9.0812e-06 - val_accuracy: 0.8601 - val_loss: 0.1257
```



```
[58] results4=model4.evaluate(x_test,y_test)
```

782/782 — 2s 2ms/step - accuracy: 0.8533 - loss: 0.6457



#### 4) Instead of relu, consider employing the tanh activation, which was well-liked in the early days of neural networks.

This code creates a neural network model (model 4) that is activated by the tanh function and has two hidden layers, each consisting of sixteen neurons. It uses Keras to accomplish this. The model is assembled using the Adam optimizer, and the loss function is the mean squared error (MSE). It is trained on a portion of the training data over 20 epochs with a batch size of 512, and its performance is monitored with validation data. The accuracy and loss statistics during training and validation epochs are shown using Matplotlib. As the model is trained, its accuracy increases. much, up to 99% or such. The model shows that it can correctly categorize the data when evaluated on a test dataset, with a test accuracy of roughly 88.0%.

```

Epoch 1/20
30/30 — 4s 85ms/step - accuracy: 0.6873 - loss: 0.5972 - val_accuracy: 0.8582 - val_loss: 0.3737
Epoch 2/20
30/30 — 3s 24ms/step - accuracy: 0.8986 - loss: 0.2988 - val_accuracy: 0.8866 - val_loss: 0.2848
Epoch 3/20
30/30 — 1s 27ms/step - accuracy: 0.9488 - loss: 0.1893 - val_accuracy: 0.8878 - val_loss: 0.2727
Epoch 4/20
30/30 — 1s 20ms/step - accuracy: 0.9603 - loss: 0.1370 - val_accuracy: 0.8852 - val_loss: 0.2839
Epoch 5/20
30/30 — 1s 20ms/step - accuracy: 0.9763 - loss: 0.0970 - val_accuracy: 0.8808 - val_loss: 0.3055
Epoch 6/20
30/30 — 1s 21ms/step - accuracy: 0.9852 - loss: 0.0723 - val_accuracy: 0.8793 - val_loss: 0.3294
Epoch 7/20
30/30 — 1s 23ms/step - accuracy: 0.9926 - loss: 0.0496 - val_accuracy: 0.8747 - val_loss: 0.3619
Epoch 8/20
30/30 — 1s 23ms/step - accuracy: 0.9959 - loss: 0.0374 - val_accuracy: 0.8716 - val_loss: 0.3907
Epoch 9/20
30/30 — 1s 21ms/step - accuracy: 0.9968 - loss: 0.0268 - val_accuracy: 0.8716 - val_loss: 0.4173
Epoch 10/20
30/30 — 1s 21ms/step - accuracy: 0.9989 - loss: 0.0188 - val_accuracy: 0.8701 - val_loss: 0.4440
Epoch 11/20
30/30 — 1s 20ms/step - accuracy: 0.9995 - loss: 0.0146 - val_accuracy: 0.8675 - val_loss: 0.4678
Epoch 12/20
30/30 — 2s 30ms/step - accuracy: 0.9997 - loss: 0.0112 - val_accuracy: 0.8686 - val_loss: 0.4912
Epoch 13/20
30/30 — 1s 29ms/step - accuracy: 0.9996 - loss: 0.0087 - val_accuracy: 0.8677 - val_loss: 0.5101
Epoch 14/20
30/30 — 1s 21ms/step - accuracy: 0.9998 - loss: 0.0066 - val_accuracy: 0.8666 - val_loss: 0.5270
Epoch 15/20
30/30 — 1s 20ms/step - accuracy: 1.0000 - loss: 0.0053 - val_accuracy: 0.8661 - val_loss: 0.5434
Epoch 16/20
30/30 — 1s 19ms/step - accuracy: 1.0000 - loss: 0.0044 - val_accuracy: 0.8664 - val_loss: 0.5579
Epoch 17/20
30/30 — 1s 21ms/step - accuracy: 1.0000 - loss: 0.0039 - val_accuracy: 0.8653 - val_loss: 0.5730
Epoch 18/20
30/30 — 1s 20ms/step - accuracy: 1.0000 - loss: 0.0034 - val_accuracy: 0.8665 - val_loss: 0.5856
Epoch 19/20
30/30 — 1s 23ms/step - accuracy: 1.0000 - loss: 0.0029 - val_accuracy: 0.8659 - val_loss: 0.5972
Epoch 20/20
30/30 — 1s 21ms/step - accuracy: 1.0000 - loss: 0.0026 - val_accuracy: 0.8658 - val_loss: 0.6081

```



```
[62] results4=model4.evaluate(x_test,y_test)
```

```

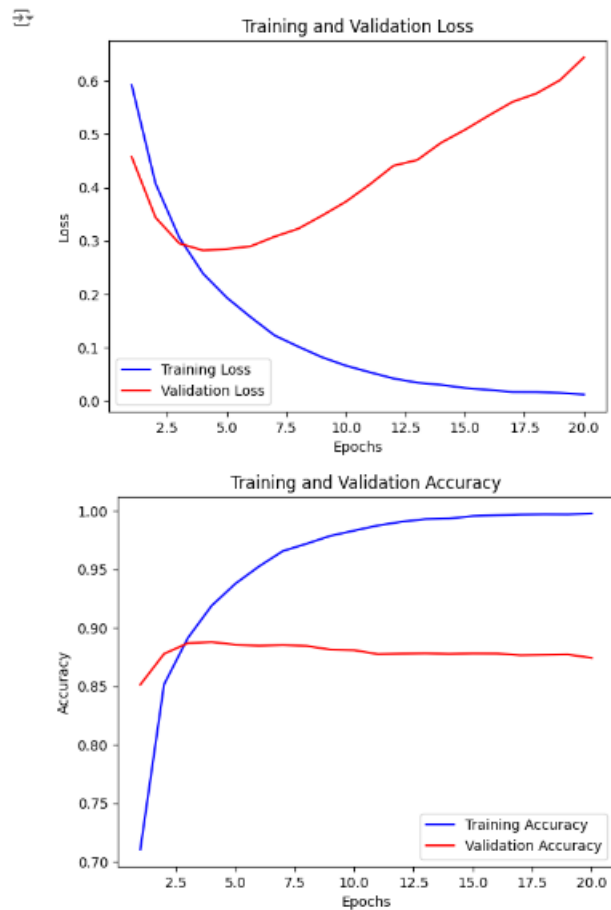
782/782 — 3s 3ms/step - accuracy: 0.8520 - loss: 0.6754

```

### 5) Make use of any technique we covered in class to improve your model's validation performance, such as dropout and regularization.

This code builds a neural network model (model 5) with a dropout layer to prevent overfitting and two hidden layers (each with 16 neurons) that are activated by ReLU. Keras is used to create the model. The binary cross-entropy loss function and the Adam optimizer are utilized in the model compilation process for binary classification tasks. It is trained on a portion of the training data across 20 epochs with a batch size of 512, and it verifies on a separate validation dataset. Throughout the training phase, accuracy has grown dramatically while loss has decreased throughout the course of the epochs. The model's test accuracy of about 88.0% indicates that it is successful in classifying the input data when tested on a test dataset.

```
Epoch 1/20  
30/30 ----- 6s 114ms/step - accuracy: 0.6429 - loss: 0.6408 - val_accuracy: 0.8513 - val_loss: 0.4573  
Epoch 2/20  
30/30 ----- 1s 20ms/step - accuracy: 0.8441 - loss: 0.4335 - val_accuracy: 0.8777 - val_loss: 0.3439  
Epoch 3/20  
30/30 ----- 1s 21ms/step - accuracy: 0.8879 - loss: 0.3192 - val_accuracy: 0.8868 - val_loss: 0.2952  
Epoch 4/20  
30/30 ----- 1s 23ms/step - accuracy: 0.9168 - loss: 0.2453 - val_accuracy: 0.8878 - val_loss: 0.2820  
Epoch 5/20  
30/30 ----- 2s 31ms/step - accuracy: 0.9372 - loss: 0.1946 - val_accuracy: 0.8855 - val_loss: 0.2846  
Epoch 6/20  
30/30 ----- 1s 27ms/step - accuracy: 0.9542 - loss: 0.1583 - val_accuracy: 0.8846 - val_loss: 0.2898  
Epoch 7/20  
30/30 ----- 1s 25ms/step - accuracy: 0.9659 - loss: 0.1240 - val_accuracy: 0.8853 - val_loss: 0.3078  
Epoch 8/20  
30/30 ----- 1s 22ms/step - accuracy: 0.9733 - loss: 0.0999 - val_accuracy: 0.8844 - val_loss: 0.3227  
Epoch 9/20  
30/30 ----- 1s 22ms/step - accuracy: 0.9785 - loss: 0.0850 - val_accuracy: 0.8813 - val_loss: 0.3473  
Epoch 10/20  
30/30 ----- 1s 20ms/step - accuracy: 0.9836 - loss: 0.0662 - val_accuracy: 0.8808 - val_loss: 0.3734  
Epoch 11/20  
30/30 ----- 1s 21ms/step - accuracy: 0.9881 - loss: 0.0523 - val_accuracy: 0.8774 - val_loss: 0.4056  
Epoch 12/20  
30/30 ----- 1s 20ms/step - accuracy: 0.9909 - loss: 0.0436 - val_accuracy: 0.8778 - val_loss: 0.4408  
Epoch 13/20  
30/30 ----- 1s 21ms/step - accuracy: 0.9930 - loss: 0.0351 - val_accuracy: 0.8781 - val_loss: 0.4516  
Epoch 14/20  
30/30 ----- 1s 24ms/step - accuracy: 0.9944 - loss: 0.0288 - val_accuracy: 0.8776 - val_loss: 0.4838  
Epoch 15/20  
30/30 ----- 1s 25ms/step - accuracy: 0.9950 - loss: 0.0249 - val_accuracy: 0.8780 - val_loss: 0.5081  
Epoch 16/20  
30/30 ----- 1s 24ms/step - accuracy: 0.9959 - loss: 0.0217 - val_accuracy: 0.8779 - val_loss: 0.5348  
Epoch 17/20  
30/30 ----- 1s 20ms/step - accuracy: 0.9973 - loss: 0.0164 - val_accuracy: 0.8765 - val_loss: 0.5605  
Epoch 18/20  
30/30 ----- 1s 24ms/step - accuracy: 0.9970 - loss: 0.0165 - val_accuracy: 0.8768 - val_loss: 0.5761  
Epoch 19/20  
30/30 ----- 1s 28ms/step - accuracy: 0.9968 - loss: 0.0156 - val_accuracy: 0.8771 - val_loss: 0.6013  
Epoch 20/20  
30/30 ----- 2s 42ms/step - accuracy: 0.9979 - loss: 0.0121 - val_accuracy: 0.8742 - val_loss: 0.6437
```



```
[66] results=model15.evaluate(x_test,y_test)
```

782/782 — 2s 2ms/step - accuracy: 0.8582 - loss: 0.7189

### Comparative Analysis of IMDB Sentiment Analysis Results from Hyperparameter Tuning.

Model no.	technique	dropout	Train accuracy	Validation accuracy	Test accuracy
1	In class method	no	100	86.7	85.7
2	Adding extra hidden layer	no	100	86.77	85.52
3	Increasing to 32 hidden layer units	no	100	86.74	85.67
4	Doing mse	no	100	86.01	85.33
4	Using tanh activation	no	100	86.58	85.2
5	Using regularization techniques	yes	99.79	87.42	85.82



**Conclusion:**

The model performed reasonably well in classifying the test data, with an accuracy of 85.82% and a loss of 0.7189. According to the accuracy metric, the model can accurately predict 86% of the situations. The loss value, which calculates the discrepancy between expected and actual results, is somewhat high, though, suggesting that the model still needs to be fine-tune. This difference between loss and accuracy could indicate that the data is either under- or overfitted by the model. Its performance may be enhanced by additional methods like regularization, hyperparameter adjustment, or experimenting with different model designs. The findings show a good beginning, but additional testing on a validation set is advised to make sure the model generalizes effectively. Deploying this approach for real-world applications would necessitate ongoing monitoring to evaluate performance on unknown data and make additional modifications as needed.